

Agentic Goal-Conditioned Workload Synthesis for Dynamic Power Characterization

Aman Desai

University of California, Santa Barbara
Santa Barbara, CA, USA
amandesai@ucsb.edu

Zishen Wan

Columbia University
New York, NY, USA
zw3306@columbia.edu

Abstract—Temporal power characterization requires executable workloads that produce specified switching-activity profiles while completing required work. Synthesizing these workloads is difficult because operation latency, contention, and persistent state couple activity across time. We present an agentic approach to goal-conditioned workload synthesis that combines hardware interface descriptions with interval-level simulation feedback. A large language model (LLM) proposes and revises complete workloads, while a fixed evaluator checks functional correctness, work completion, deadlines, and activity tolerance in every interval. Search uses register-transfer-level (RTL) switching activity; selected workloads undergo separate mapped-gate power assessment. Across five designs, Gemini 3.8 Flash satisfies the activity specification in 387 of 400 runs on time-varying targets (96.8%) within 128 proposals per run. Across the four non-Mesh designs, it reduces the area under the activity-error curve by 85.9-89.8% relative to the best evaluated classical baseline per design. The resulting workloads provide executable stimuli for controlled activity experiments and subsequent power analysis.

Index Terms—workload generation, dynamic power, hardware agents, temporal activity, design automation

I. INTRODUCTION

A hardware workload determines not only what a circuit computes, but also which parts of the circuit switch and when they switch. These activity patterns affect dynamic power, so an average alone cannot distinguish a short burst from sustained demand or repeated changes in activity. Temporal stressmark generation already uses controlled changes in demand to study power-delivery behavior [1], [2]. For a designer who wishes to examine a particular temporal condition, power analysis is only part of the task. The designer also needs a workload that produces that condition.

Consider a packet network that must alternate between busy and quiet intervals while delivering its scheduled traffic. Choosing packet release times does not fully determine when switching occurs. Contention and backpressure can delay delivery, allowing activity from one interval to continue into the next. Similar dependencies arise when instruction latency shifts later processor work. The next proposal must account for both available operations and their measured timing effects. Can an LLM use supplied interface descriptions and interval-level errors to find better-matching workloads than classical search within the same proposal budget?

We investigate this question across five hardware designs. An LLM proposes complete workloads from interface de-

scriptions and revises its proposals using simulation feedback. Random, evolutionary, and surrogate-guided policies use the same workload languages and proposal budgets.

During search, RTL switching activity guides workload generation without requiring gate-level simulation for every proposal. These transition counts serve as a search proxy, not a direct power estimate. To assess power, we replay selected and reference workloads through the synthesized designs' interfaces. The resulting workloads can be inspected, executed, and reused.

II. PROBLEM AND MEASUREMENT

Given a fixed design, required work, and eight interval targets, we search for a workload that completes the work correctly and matches those targets. We evaluate workload w by simulating it on design d (Figure 1):

$$\tau = \text{Simulate}(d, w), \quad a_k = \frac{\text{Toggles}(\tau, I_k)}{C_k}. \quad (1)$$

Here τ is the signal trace and I_k contains C_k clock cycles. Toggles counts known 0-to-1 and 1-to-0 bit transitions in the measured RTL scope, excluding clocks. Thus a_k is transitions per cycle, in the same units as the requested activity t_k .

A. Tracking error and acceptance

The observation period has $K = 8$ consecutive intervals, or bins, with fixed boundaries per design and lengths differing by at most one cycle. To interpret errors on a consistent scale, we divide by a design-specific activity span $s_d > 0$: the 95th minus fifth percentile of valid calibration interval rates. This scale stays unchanged across targets and proposals. Signed errors describe each interval, while root-mean-square (RMS) and maximum error summarize the profile:

$$e_k = (a_k - t_k)/s_d, \quad k = 1, \dots, K, \\ E_{\text{RMS}} = \sqrt{\frac{1}{K} \sum_{k=1}^K e_k^2}, \quad E_{\text{max}} = \max_k |e_k|. \quad (2)$$

Acceptance requires correct execution, completion of the required work on time, and $E_{\text{max}} \leq \epsilon = 0.05$. Requiring every interval to meet the tolerance prevents accurate matches in some intervals from compensating for a large error in another. The tolerance is relative to the calibration span, not

ACTIVITY-GUIDED SEARCH




Fig. 1. Simulation feedback guides workload proposals. After search, the selected workload and a separately constructed reference are replayed on the synthesized design under the same conditions. Estimated power assesses the result but does not guide search.

a percentage of target activity or power. RMS error measures progress toward a match.

We track the lowest RMS error among valid workloads found so far. Its trapezoidal area under the curve (AUC) measures progress over $N = 128$ proposal slots, not elapsed time. Invalid, missing, and duplicate proposals count toward the budget. After a successful batch, search stops and its best error is carried forward to N . Unsuccessful runs remain in the summaries. Extended methods give the full accounting convention.

III. GOAL-CONDITIONED WORKLOAD SYNTHESIS

A. An executable feedback loop

The model proposes a complete workload using the design’s available operations. An adapter checks its structure and interface constraints, then translates it into simulator inputs. Simulation reveals when activity actually occurs. For valid workloads, the evaluator reports which intervals have too much or too little activity. Otherwise, it reports the rejection reason.

The next request supplies the workload format, interface description, target and scale, and a compact history of outcomes. Errors identify what needs improvement, while the interface description identifies possible edits. Those edits can affect several intervals: Ibex polling waits still switch, and earlier RedMule jobs can delay later jobs beyond their release times. Simulation captures these effects rather than equating workload phases with measurement intervals. The model cannot alter the RTL or measurement rules.

B. Baselines with temporal control

Mesh baselines retain sink period 8 and pauses from 0 to 3. The LLM schema permits wider settings. **Phase-random** samples workload phases, their timing and operations. **Phase-GA** selects workloads using measured RMS error, then changes their structure and timing to produce new candidates. **Phase-model** fits ridge regression after eight valid measurements to

TABLE I
EVALUATED HARDWARE INTERFACES AND CONTROLLABLE WORKLOAD PARAMETERS.

Design	Workload control
OpenTitan AES	Encryption grouping and waits, fixed key/data
AXI DMA	Fixed-size copy grouping and pacing, no injected backpressure
Ibex	Instruction bodies, initial state, segment weights and releases
BaseJump mesh	Packet timing, routes, data patterns and sink behavior
RedMule	FP16 GEMM size, operands, job counts and releases

predict interval errors. It screens up to 32 edits using those predictions, also includes random candidates, and favors improvements to the worst interval. Screening uses no additional simulations, but submitted candidates consume proposal slots. Phase-model changes selection as well as prediction, rather than testing a predictor alone.

C. Gate-level replay

From each run, we select the valid workload with the smallest worst-interval error, breaking ties by RMS error and earliest slot. Unsuccessful runs contribute their best valid attempt, or no measurement if none is valid. Selection precedes power estimation so that power cannot influence the choice.

Because an activity target does not uniquely determine power, we compare selected workloads with separately constructed references. Both run on the same synthesized design with the same library, 10 ns clock, and intervals. OpenSTA estimates internal and switching power from gate activity and Liberty characterization, excluding leakage [3].

Normalized RMS power error (NRMSE) is the eight-interval RMS candidate-minus-reference error divided by the reference’s duration-weighted mean dynamic power. Profiles

TABLE II

MEAN AUC OF BEST-SO-FAR NORMALIZED RMS ERROR AND SOLVES OUT OF 80 NONFLAT RUNS PER DESIGN. BOLD MARKS THE LOWEST MEAN AUC, NOT SIGNIFICANCE. GEMINI 3.8 DENOTES GEMINI 3.8 FLASH.

Design	Phase-random		Phase-GA		Phase-model		Flash-Lite		Gemini 3.8	
	AUC ↓	Solves	AUC ↓	Solves	AUC ↓	Solves	AUC ↓	Solves	AUC ↓	Solves
AES	18.43	1	19.60	2	12.13	33	17.38	21	1.23	80
DMA	32.80	0	33.11	1	30.17	0	40.52	1	3.72	78
Ibex	44.86	0	41.75	1	46.91	0	51.31	4	4.86	80
Mesh	36.50	0	31.41	0	32.76	0	24.62	20	2.24	80
RedMule	45.72	0	37.83	0	46.57	0	40.63	0	5.33	69




Fig. 2. Both LLMs match these Mesh targets at seed 9100. Steps show 1,025-cycle average activity divided by the fixed scale, not within-interval changes. Shading marks the ± 0.05 tolerance. Counts include initialization and full batches. Selection minimizes valid $E_{\max}$ with earliest-slot ties. Classical runs use 128 slots. These are individual examples.

are neither shifted nor independently rescaled. Each reference provides one power realization of the activity target, not a unique conversion of that target into watts.

Gate-level replay uses zero delay and omits routed parasitics and delay-induced glitches. Since OpenSTA limits switching rates based on modeled signal transition times, short bursts can exceed its limit even when the whole-run average does not. Interval-level power estimates can therefore differ from the whole-run estimate. For Ibex and RedMule, we verify that clipping explains each discrepancy without changing the reported power. Functional replay, alignment, finite power and annotation checks remain required.

IV. EVALUATION DESIGN

Each design has eight time-varying (nonflat) targets and a near-flat control. Ten seeds give 80 nonflat runs per design, or 400 across 40 tasks per policy. For each target and seed, policies share two initial workloads and 128 proposal slots. Both candidates in each LLM response count, even if the first matches. Observation periods span 6,774 to 262,144 cycles. Work requirements prevent quiet matches obtained by omitting work. Targets include steps, bursts, quiet intervals, alternation, ramps and irregular levels. Calibration measurements set their amplitudes. AES and DMA preserve mean requested activity across shapes, so bursts request more activity in fewer intervals.

Gemini 3.5 Flash-Lite and Gemini 3.8 Flash use MEDIUM reasoning, temperature 0.7 and top-p 0.95, with output-token allowances of 8,192 and 65,536 respectively.

RTL simulation uses Verilator, except for DMA’s Icarus flow. Yosys/Slang synthesizes to Sky130 HD at typical 25°C, 1.8 V. Released CHIA bindings expose tool stages, while scores use the frozen benchmark evaluator.

For each design, we average paired AUC differences across eight targets per seed. Ten averages support bootstrap intervals and exact sign-flip tests. We adjust p-values across each LLM’s 15 baseline comparisons to limit the chance of any false-positive finding in that group to 5%. Evaluation uses an already-observed bank, not fresh held-out tasks.

V. RESULTS AND GATE-LEVEL POWER ASSESSMENT

A. Completed Search Comparisons

Gemini 3.8 Flash matches 387 of 400 nonflat runs, compared with 46 for Flash-Lite and 1, 4, and 33 for Phase-random, Phase-GA, and Phase-model (Table II). It also finds better activity matches earlier in search, reducing mean error-curve AUC by 85.9% to 92.9% relative to the lowest-AUC classical baseline on each design. All 15 baseline comparisons favor Gemini 3.8 and remain significant after adjustment for multiple testing. Flash-Lite improves on the classical policies on Mesh but has higher AUC than all of them on DMA. Gemini 3.8 also has a larger token allowance.

TABLE III
MEAN ERROR AGAINST THE REFERENCE WORKLOAD’S ESTIMATED POWER PROFILE (NRMSE, % OF REFERENCE MEAN). CLASSICAL SELECTS PHASE-MODEL FOR AES/DMA AND PHASE-GA OTHERWISE BY LOWEST MEAN ERROR. COLUMNS SHARE n MEASURED PAIRS.

Design	Gemini 3.8	Classical	n
AES	0.956	8.675	80
DMA	0.076	1.031	80
Ibex	14.481	36.848	80
Mesh	0.670	1.559	80
RedMulE	6.143	24.364	80

Across nonflat runs, Gemini 3.8 uses 6,008 charged proposals, compared with 49,290 to 51,088 for the classical policies. Counts include initialization and unsuccessful runs. Both LLMs match the targets in Figure 2, but Gemini 3.8 reaches tolerance earlier without always having smaller final error. The illustrated Flash-Lite workload alternates 35-packet neighbor phases and 320-packet hotspot phases. All 1,420 packets complete and every interval meets tolerance. Its 900-cycle release phases need not align with measurement bins.

Only Gemini 3.8 solves nonflat RedMulE requests. Across all designs, controls yield 50/50 Gemini matches and 23/50 Flash-Lite matches, separate from nonflat totals. We replay the Mesh workloads using the baselines’ sink settings, with period 8 and pauses capped at 3. All 80 Gemini and 20 Flash-Lite nonflat matches remain within tolerance, while one Flash-Lite control no longer matches. Mean reference-power error changes by less than 0.001 percentage point for either model.

B. Mapped-Gate Assessment

Without gate-power feedback during search, Gemini 3.8 also lowers mean error against reference workloads’ estimated power on every design (Table III). All 40 nonflat references meet the activity tolerance. Each design’s power comparison includes all 80 nonflat runs per policy. Power agreement is a comparison with those workloads, not a separate power-target success rate. Of 400 nonflat Gemini profiles, 399 beat the reference’s constant approximation, its eight-bin mean. Nearly flat references can have small errors without close temporal tracking. DMA’s peak-to-trough variation is 3.0% to 4.3% of mean power.

VI. RELATED WORK

Stressmark generators seek demanding power and switching conditions [4], [5]. SAGA uses predictions to reduce measurements during genetic search [6], while SAT-based methods construct functional switching sequences [7]. Our baselines use evolutionary and surrogate search, not SAT. DiffPower provides differentiable power analysis [8], whereas we estimate power after search.

Temporal methods move beyond high average demand. Di/dt stressmarks construct changes in current [1], [2], while PaTGen combines phases and adjusts their timing to reproduce performance behavior [9]. We study requested activity profiles across different RTL workload interfaces. Related generation

loops pursue other objectives: TheHuzz and ChipFuzzer seek bugs [10], [11], and AgentDSE changes architecture rather than stimulus on fixed hardware [12]. CHIA provides composable infrastructure for such experiments [13].

VII. DISCUSSION AND CONCLUSION

Gemini 3.8 Flash meets the activity specification in 387 of 400 nonflat runs and lowers reference-power error on all five designs. These comparisons evaluate complete policy configurations, not model choice alone. The Mesh replay tests selected workloads but does not establish search performance under restricted settings. Eight intervals leave finer timing unresolved. The resulting workloads can be inspected and replayed to study temporal activity and assess estimated dynamic power. Code, workloads, and results are available at <https://github.com/amanpdesai/agecws>.

ACKNOWLEDGMENT

AI tools supported proofreading and editing of the manuscript. AI-assisted programming tools were also used to develop the accompanying software.

REFERENCES

- [1] Y. Kim and L. K. John, “Automated di/dt stressmark generation for microprocessor power delivery networks,” in *Proc. ISLPED*, 2011.
- [2] Z. Jiang, J. Garrigus, A. Seigler, E. Syed, Y.-L. Huang, M. Sadi, T. Rahal-Arabi, and L. K. John, “Exploration of LLM workload reliability based on di/dt effects and voltage droops,” in *Proc. HPCA*, 2026. [Online]. Available: https://lca.ece.utexas.edu/pubs/jiang_hPCA26.pdf
- [3] The OpenROAD Project, “OpenSTA: Static timing analysis,” <https://github.com/The-OpenROAD-Project/OpenSTA>, accessed: 2026.
- [4] A. M. Joshi, L. Eeckhout, L. K. John, and C. Isen, “Automated microprocessor stressmark generation,” in *Proc. HPCA*, 2008.
- [5] Z. Hadjilambrou, S. Das, P. N. Whatmough, D. Bull, and Y. Sazeides, “GeST: An automatic framework for generating CPU stress-tests,” in *Proc. ISPASS*, 2019. [Online]. Available: <https://github.com/toolsForUarch/GeST>
- [6] P. Chatzimiltis, G. Antoniou, H. Volos, and Y. Sazeides, “SAGA: A surrogate assisted genetic algorithm for fast CPU power virus generation,” in *Proc. ISPASS*, 2025.
- [7] N. Deligiannis, R. Cantoro, T. Faller, T. Paxian, B. Becker, and M. S. Reorda, “Effective SAT-based solutions for generating functional sequences maximizing the sustained switching activity in a pipelined processor,” in *Proc. ATS*, 2021.
- [8] I. Jacobson, Z. Zhao, R. Mehrotra, G. Zhou, V. Rashingkar, and Y. Chen, “DiffPower: GPU-accelerated differentiable switching power analysis and optimization,” *arXiv preprint arXiv:2608.03778*, 2026.
- [9] H. Yin, W. Lin, H. Huang, X. Luo, H. Zhong, and K. Li, “PaTGen: Temporal similarity-driven proxy benchmark generation method for cloud workloads,” *ACM Transactions on Architecture and Code Optimization*, vol. 23, no. 2, 2026.
- [10] R. Kande, A. Crump, G. Persyn, P. Jauernig, A.-R. Sadeghi, A. Tyagi, and J. Rajendran, “TheHuzz: Instruction fuzzing of processors using golden-reference models for finding software-exploitable vulnerabilities,” in *Proc. USENIX Security*, 2022. [Online]. Available: <https://www.usenix.org/conference/usenixsecurity22/presentation/kande>
- [11] K. Wang, C. Li, K. Chang, S. Cai, Y. Han, and Y. Wang, “When fuzzing meets understanding: LLM-driven semantic test generation for RTL verification,” *arXiv preprint arXiv:2607.10340*, 2026.
- [12] C. Wang, J. C. Shi, D. Kong, D. S. Boning, Z. Wan, Y. Du, and V. J. Reddi, “AgentDSE: Reasoning-augmented architectural design space exploration,” *arXiv preprint arXiv:2606.21836*, 2026.
- [13] A. Cui, F. Hermida-Rivera, J. Toubes, R. Gupta, J. Fang, C. L. Zhang, E. Schwarz, J. Kim, Y. S. Shao, B. Nikolic, C. W. Fletcher, and S. Karandikar, “CHIA: An open-source framework for principled, agentic AI-driven hardware/software co-design research,” *arXiv preprint arXiv:2606.27350*, 2026.